

C++ implementation of the Metropolis-Hasting Algorithm

Nicholas Malaya

Institute for Computational Engineering and Sciences
The University of Texas at Austin

December 2nd, 2014

Motivation

Extending MC

- ▶ We discussed Monte-Carlo (MC) methods in class
- ▶ But we can do better— Markov Chain Monte-Carlo (MCMC) is an extension of MC

MCMC is ubiquitous

- ▶ statistics, physics, chemistry, biology, engineering, business, linguistics, etc.
- ▶ Especially research: Machine Learning, Uncertainty Quantification, Astrophysics, Particle Physics, Risk assessment
- ▶ Primarily used for calculating numerical approximations of multi-dimensional integrals

MCMC

What is it?

- ▶ MC = Monte-Carlo
- ▶
- ▶
- ▶ Monte-Carlo = Stochastic (i.e. random numbers)

Essentially, a method to generate random

MCMC

What is it?

- ▶ CMC = Chain Monte-Carlo
- ▶
- ▶ Chain = A series of steps
- ▶ Monte-Carlo = Stochastic (i.e. random numbers)

Essentially, a method to generate random steps

MCMC


What is it?

- ▶ MCMC = Markov Chain Monte-Carlo
- ▶ Markov = Each step only depends on previous
- ▶ Chain = A series of steps
- ▶ Monte-Carlo = Stochastic (i.e. random numbers)

Essentially, a method to generate random steps correlated with the previous step


Enter Metropolis-Hastings

Intuitive Picture: Sampling a Gaussian Distribution




Enter Metropolis-Hastings

Intuitive Picture: Sampling a Gaussian Distribution



Enter Metropolis-Hastings

Intuitive Picture: Sampling a Gaussian Distribution



Enter Metropolis-Hastings

What is it?

- ▶ Metropolis-Hastings is an MCMC method for generating these samples.

Algorithm

- ▶ Start at current step (ϕ_1)
- ▶ Draw proposal: $q(\tilde{\phi}_2|\phi_1) = N(\tilde{\phi}_2|\phi_1, \nu^2)$
- ▶ Draw from uniform: $u \sim U(0, 1)$
- ▶ if: $u < \frac{\pi_n(\tilde{\phi}_2)}{\pi_n(\phi_1)}$ (π_n is bayesian posterior, etc.)
 - ▶ $\phi_2 = \tilde{\phi}_2$
- ▶ else: $\phi_2 = \phi_1$

Enter Metropolis-Hastings

What is it?

- ▶ Metropolis-Hastings is an MCMC method for generating these samples.

Algorithm

- ▶ Start at current step (ϕ_1)
- ▶ Draw proposal: $q(\tilde{\phi}_2|\phi_1) = N(\tilde{\phi}_2|\phi_1, \nu^2)$
- ▶ Draw from uniform: $u \sim U(0, 1)$
- ▶ if: $u < \frac{\pi_n(\tilde{\phi}_2)}{\pi_n(\phi_1)}$ (π_n is bayesian posterior, etc.)
 - ▶ $\phi_2 = \tilde{\phi}_2$
- ▶ else: $\phi_2 = \phi_1$

Question: How do we choose ν ?

Project Status

Library Snapshot



- ▶ C++, threaded with OpenMP
- ▶ Makefile build system support for:
 - ▶ Asus laptop
 - ▶ TACC's Lonestar supercomputer
- ▶ Git revision control
- ▶ CLOC: 254 SLOC C++

Capabilities

- ▶ Gamma, Beta, Normal Statistical Distributions

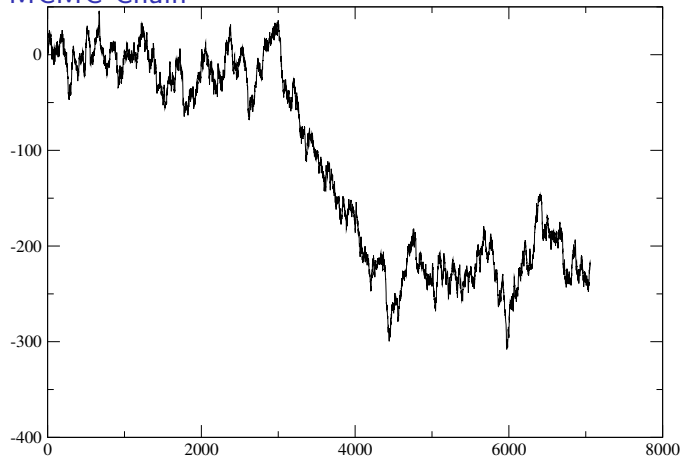
Metropolis-Hastings Sampling in C++

Sampling a Gamma Distribution



Burn-in

MCMC Chain



OpenMP Threading

Embarassingly Parallel

- ▶ Each chain is completely independent
- ▶ Thus, each thread can run completely independent, aside from ALL_REDUCE at end
- ▶ As a result, we anticipate nearly-perfect scalability.

Threading the application

- ▶ compile with: `g++ test.cpp -fopenmp -lpthread`
- ▶ `OMP_GET_NUM_PROCS != OMP_GET_MAX_THREADS`
 - ▶ Why might that be?

Conclusions

Summary

- ▶ MCMC is a critical algorithm across a wide variety of industries and STEM fields
- ▶ We have demonstrated a simple C++ implementation
- ▶ This implementation should scale well on a single lonestar node

Thank you!

- ▶ nick@ices.utexas.edu for code
- ▶ Have a Markovian Day!